

A Coq Library for Verification of Concurrent Programs

Reynald Affeldt^{a,1} Naoki Kobayashi^{b,2}

^a *Research Center for Information Security, National Institute of Advanced Industrial Science and Technology, Tokyo, Japan*

^b *Graduate School of Information Sciences, Tohoku University, Sendai, Japan*

Abstract

Thanks to recent advances, modern proof assistants now enable verification of realistic sequential programs. However, regarding the concurrency paradigm, previous work essentially focused on formalization of abstract systems, such as pure concurrent calculi, which are too minimal to be realistic. In this paper, we propose a library that enables verification of realistic concurrent programs in the Coq proof assistant. Our approach is based on an extension of the π -calculus whose encoding enables such programs to be modeled conveniently. This encoding is coupled with a specification language akin to spatial logics, including in particular a notion of fairness, which is important to write satisfactory specifications for realistic concurrent programs. In order to facilitate formal proof, we propose a collection of lemmas that can be reused in the context of different verifications. Among these lemmas, the most effective for simplifying the proof task take advantage of confluence properties. In order to evaluate feasibility of verification of concurrent programs using this library, we perform verification for a non-trivial application.

Keywords: Proof assistant, Coq, concurrent programs, pi-calculus

1 Introduction

Concurrent programs are ubiquitous: multi-threaded programs in network servers, distributed programs for database applications, etc. In order to guarantee their correctness and security properties, it is important to verify them formally. The main difficulty in formally verifying concurrent programs is the size of their state space. The latter can be very large (because of non-determinism) and even infinite (for non-terminating applications, such as reactive systems).

Proof assistants and model checkers can be regarded as complementary tools for formal verification. Model checkers are fully automated but can only handle finite state space systems (without appropriate abstraction techniques). Proof assistants

¹ Email: reynald.affeldt@aist.go.jp

² Email: koba@ecei.tohoku.ac.jp

are interactive but they can handle infinite state space systems directly, using inductive reasoning. In this paper, we are concerned with formal verification based on proof assistants.

Proof assistants have been successfully applied to formal verification of sequential programs. There are realistic use-cases (e.g., [2,5,27]) and tools that enable practical verification of imperative programs (e.g., [14]).

Regarding concurrency, previous work using proof assistants has focused on abstract concurrent systems rather than on realistic concurrent programs. There are many formalizations of pure concurrent calculi (e.g., [13,17,18,28,29]) and experiments with the combined use of proof assistants and model checkers for minimal concurrent languages (e.g., [22,32]). This work demonstrates the usefulness of proof assistant-based formal verification for concurrent programs. However, the formalized calculi and languages are cumbersome to verify realistic concurrent programs because of the lack of datatypes. Moreover, in view of the large proof developments in previous work, it is questionable whether such verifications can be done in practice. For these reasons, we think that formal verification of realistic concurrent programs has not yet been addressed satisfactorily.

In this paper, we introduce a library that enables verification of realistic concurrent programs using a general-purpose proof assistant, namely Coq [30]. This library consists of:

- A modeling language with attractive features for verification of realistic concurrent programs. This modeling language is based on the π -calculus [24] (a foundational language for the study of concurrent systems) but is different from encodings developed in previous work in that it allows Coq datatypes and control structures to be used. Consequently, it makes it easy to model realistic concurrent programs and to run these models using existing virtual machines and compilers.
- A specification language for realistic concurrent programs. In particular, it provides a notion of fairness that is necessary to write satisfactory specifications for realistic concurrent programs.
- A collection of lemmas in order to facilitate formal proof. The most effective lemmas are based on confluence properties. They allow for smaller formal proofs by reducing the state space that needs to be explored for the purpose of verification.

To evaluate the feasibility of verification of concurrent programs using our library, we have performed formal verification of an existing mail server.

In the rest of this paper, we explain the three parts of the library in turn (the modeling language, the specification language and the collection of lemmas) and then report on the case study. We use the syntax of Coq (version 8).

2 Modeling Language

In this section, we introduce (a Coq encoding of) a simple concurrent language that can be used to model a wide range of realistic concurrent programs. Simplicity and generality are inherited from the π -calculus, on which this modeling language is

based. Because of its minimality, the (pure) π -calculus is not well-suited to modeling of realistic concurrent programs. The main reason is that datatypes and control structures (conditionals and functions) need to be encoded by means of the concurrent primitives. Our modeling language addresses this shortcoming by extending the π -calculus with datatypes and functions, similarly to the Pict programming language [26]. We call our modeling language $\text{appl}\pi$, which stands for “applied π -calculus”³. In Sect. 2.1 and Sect. 2.2 we discuss the encoding of the syntax and the operational semantics of $\text{appl}\pi$, respectively.

2.1 Syntax Encoding

The syntax of $\text{appl}\pi$ consists of *channels* and *processes*. Intuitively, processes perform computations and exchange values with other processes through channels.

Channels are encoded by means of the functional type **chan**. Any type in **Set** can be used as a datatype for communicated values, and channels themselves can be communicated: **Axiom** $\text{chan} : \text{Set} \rightarrow \text{Set}$.

Processes are encoded by means of the inductive type **proc**. Each constructor of **proc** corresponds to a concurrent primitive of the π -calculus⁴:

```
Inductive proc : Type :=
  zeroP: proc
| inP  : forall A, chan A -> (A -> proc) -> proc
| rinP : forall A, chan A -> (A -> proc) -> proc
| outP : forall A, chan A -> A -> proc -> proc
| parP : proc -> proc -> proc
| nuP  : forall A, (chan A -> proc) -> proc.
```

Intuitively, **zeroP** represents the inert process. **inP** $c \text{ (fun } x:A \Rightarrow P)$ represents an input process: it waits for some value v of type A along the channel c and then behaves as process $((\text{fun } x:A \Rightarrow P) v)$. **outP** $c v P$ represents an output process: it sends the value v along the channel c and then behaves as process P . **parP** $P Q$ represents the parallel composition of the processes P and Q . **rinP** $c \text{ (fun } x:A \Rightarrow P)$ represents replicated input: it waits for some value v of type A along the channel c and then behaves as process **parP** (**rinP** $c \text{ (fun } x:A \Rightarrow P)$) $((\text{fun } x:A \Rightarrow P) v)$. The process **nuP** $(\text{fun } x:\text{chan } A \Rightarrow P)$ represents channel creation: it creates a new channel c' and then behaves as the process $((\text{fun } x:\text{chan } A \Rightarrow P) c')$. Processes are in the **Type** universe so that they cannot be sent as data.

This encoding allows the Coq language to be used as the functional core of $\text{appl}\pi$. This effect is achieved by *higher-order abstract syntax* (HOAS), an encoding technique used to ease the management of binders. Concretely, process continuations for input and channel creation primitives are taken to be Coq functions. Thus, one can use the Coq language to write $\text{appl}\pi$ processes. Though convenient, this feature

³ We introduce this abbreviation to avoid confusion with Abadi and Fournet’s applied π -calculus [1] which is an extension of the π -calculus to study security protocols.

⁴ The concurrent primitives of $\text{appl}\pi$ are more precisely a subset of those of the π -calculus: replication is restricted to input processes and there is no external choice. These restrictions have little impact on expressiveness, as discussed in [26].

allows for “exotic terms” [12], i.e., terms that do not correspond to any process of the π -calculus. For example, the following process is such an exotic term (c has type `chan nat`):

```
inP c (fun n => match n with 0 => zeroP | S _ => parP zeroP zeroP end)
```

We do not consider this to be an issue in this paper because we are concerned with modeling of concurrent programs rather than on adequacy of the encoding. Yet, since we are dealing with a language larger than the π -language, we have to be careful when axiomatizing not to take the properties of the latter for granted.

Our encoding can be said to be a *deep embedding* because we define the syntax as an inductive type that we use in the next section to define the operational semantics. However, the ability to integrate Coq functions gives it also the flavor of a *shallow embedding*. See [6] and [23] for definitions.

The use of dependent types guarantees that channels are used consistently according to their type. For example, `inP c (fun x:A => P)` is rejected by Coq if c has not type `chan A`. Without dependent types, we would have to introduce a sum type for values and insert explicit tagging/untagging to perform data emission and reception, what would make modeling in $\text{appl}\pi$ cumbersome. The combined use of HOAS and dependent types makes our encoding different from previous work on encoding of the π -calculus in Coq.

The following definitions are used in the rest of the paper. They represent output and input processes without continuations:

Definition `OutAtom (A:Set) (c:chan A) (v:A) := outP c v zeroP.`

Definition `InAtom (A:Set) (c:chan A) := inP c (fun x => zeroP).`

Before discussing the formal operational semantics, we illustrate the practical advantages of $\text{appl}\pi$ as a modeling language.

2.1.1 Modeling Realistic Concurrent Programs

By way of example, we show below how to model a simple client/server program.

The process below represents a simple server. It waits on the channel i for a request, here a pair of a natural number and a channel. It computes the successor of the received natural number and sends it back using the received channel:

Definition `server (i:chan (nat * (chan nat))) : proc :=`
`rinP i (fun ar => let a := fst ar in let r := snd ar in`
`OutAtom r (plus a 1)).`

Observe that it is easy to write realistic programs because our encoding provides us with the Coq language and the Coq standard library (here: `let` construct; `plus`, `fst`, `snd` functions).

The process below represents a client for the above server. It sends a request and waits for the answer of the server along a channel it has created. Eventually, it displays the server response along the channel o :

Definition `client (i:chan (nat * (chan nat))) (o:chan nat) : proc :=`

```
nuP1 (fun r => parP (OutAtom i (0,r)) (inP r (fun x => OutAtom o x))).
```

The parallel composition (`parP (server i) (client i o)`) models a simple client-server program. We discuss further modeling issues in Sect. 5.

2.1.2 Executing $\text{appl}\pi$ Models

It is possible to run $\text{appl}\pi$ models with little modification by using the extraction facility of Coq. For instance, the server above can be turned into OCaml code:

```
Coq < Recursive Extraction server.
...
(* various OCaml data structures and functions, including
a datatype for concurrent primitives and the plus function *)
...
let server i =
  RinP ((Obj.magic i), (fun ar ->
    outAtom (snd (Obj.magic ar)) (plus (fst (Obj.magic ar)) (S 0))))
```

To run that program using existing virtual machines or compilers, it is sufficient to replace the type constructors for concurrent primitives by OCaml functions with the appropriate semantics. For a sample OCaml module with such functions, see [3]. This facility can be used to run $\text{appl}\pi$ models as programs on their own. More radically, one can use $\text{appl}\pi$ not as a modeling language but as a programming language, the Coq interface providing static type checking (a la polyadic π -calculus, thanks to our use of dependent types) and OCaml providing an efficient execution environment for formally verified programs.

2.2 Operational Semantics Encoding

The operational semantics of $\text{appl}\pi$ is a relation between processes, which defines what it means for a process to execute actions such as data emission/reception and channel creation. Similarly to the syntax, the operational semantics is borrowed from the π -calculus. More precisely, it is a non-standard labeled transition semantics. Before explaining the encoding, we justify the need for a non-standard semantics.

Our use of HOAS makes it difficult to encode the standard semantics of the π -calculus. The difficulty comes from the fact that ν -bound channels and conditionals are handled at the meta-level in our syntax encoding (respectively by Coq variables and Coq case analysis). For illustration, let us consider the following $\text{appl}\pi$ process:

```
nuP (fun x => parP (inP x (fun _ => OutAtom x v)) (OutAtom x v))
```

Using a standard semantics, we would expect it to reduce by communication along channel x to the process `nuP (fun x => OutAtom x v)`. It is difficult to write in Coq a rule to perform such reductions because the processes that are reduced are inside a meta-level λ -abstraction. Honsell et al. [18] solve this problem in their HOAS encoding of the (pure) π -calculus by introducing “freshness” predicates. For example, their encoding of the standard rule for channel creation requires an additional

predicate to check occurrence of a channel in a process (predicate `notin` in the rule `fRES` in [18]). However, this solution is not directly applicable to `appl π` because conditionals are represented by Coq case analysis (whereas they are represented by type constructors in [18]).

Our solution is to distinguish between channels already created and channels to be created. For this purpose, instead of considering sole processes, we consider *states*, i.e., pairs of a process with the list of the channels created so far. In a state, channels already created appear in the list and the channels to be created appear as ν -bound channel in the process. (In comparison, both kinds of channels are represented by ν -bound channels in standard semantics.) We note `L#P` the state composed of the list `L` and the process `P`, `nilC` the empty list, and `&` the addition of an element to a list. The `appl π` process above is rewritten into the state:

```
nilC#nuP (fun x => parP (inP x (fun _ => OutAtom x v)) (OutAtom x v))
```

Using our non-standard semantics, it first creates a new channel `x'` to replace `x`:

```
x'&nilC # parP (inP x' (fun _ => OutAtom x' v)) (OutAtom x' v)
```

and then reduces by communication along channel `x'`:

```
x'&nilC # OutAtom x' v
```

Concretely, the operational semantics is encoded by means of two inductive predicates `Trans` and `Redwith`. `Trans P l Q` means that process `P` reduces to process `Q` by performing the elementary action `l` (of type `TrLabel`, representing either data emission, data reception, channel creation or communication). The formal definition of the `Trans` predicate is similar to standard labeled transition semantics except for the rule for channel creation:

```
Inductive Trans : proc -> TrLabel -> proc -> Prop :=
...
| tr_new : forall A (C:chan A -> proc) (c:chan A),
  Trans (nuP C) (NewL c) (C c).
```

`Redwith S l S'` means that state `S` reduces to state `S'` by performing a communication or a channel creation (action of type `RedLabel`). In particular, it captures what it means for a channel to be new (or “fresh”): simply that it does not appear in the list of channels created so far.

```
Inductive Redwith : state -> RedLabel -> state -> Prop :=
...
| red_new : forall L P Q A (c:chan A),
  Trans P (NewL c) Q -> fresh c L -> Redwith (L # P) (New c) (c&L # Q).
```

In the following, when `Redwith S l S'` is true for some `l`, we write `Red S S'`. We also write `Reds` for the reflexive, transitive closure of `Red`.

3 Specification Language

Specification of concurrent programs deals with questions such as reachability of desirable states. There are several specification languages (or logics) designed for that purpose, such as spatial logics [7] or Dam’s π - μ -calculus [11]. The specification language provided in our library is based on Cardelli and Gordon’s spatial logic [7], because we found it expressive enough for our purpose.

Concerning temporal formulas, an important issue developed in our specification language is formalization of strong fairness. Intuitively, strong fairness is a system property enjoyed by execution environments in which communications that can execute infinitely often are eventually scheduled for execution. It is an important assumption without which we cannot write satisfactory specifications for realistic concurrent programs. For example, let us consider the following program:

```
parP (parP (OutAtom d v) (inP d (fun _ => OutAtom e v)))
      (parP (OutAtom c v) (rinP c (fun _ => OutAtom c v)))
```

A property that one might want to check is that the process `OutAtom e v` is eventually revealed. However, without the fairness assumption, this property does not even hold. We show in Sect. 3.1 how we encode the fairness assumption and in Sect. 3.2 we give the semantics of the formulas of our specification language.

3.1 Encoding of the Fairness Assumption

Fairness is expressed by means of quantifications over runs of concurrent programs. We first explain how we encode runs.

Encoding of Runs

A run intuitively consists of a maximal sequence of successive reductions. A *state sequence* is an indexed set of optional states. A *stable* state is a state that cannot evolve anymore.

Definition `stateSeq` : Type := nat -> optionT state.

Definition `Stable` (P:state) : Prop := ~(exists Q, Red P Q).

A *reduction sequence* is a state sequence such that each state is obtained by a reduction of its predecessor:

Definition `isRedSeq` (PS:stateSeq) : Prop := forall n,
 (forall P, PS n = SomeT _ P ->
 (exists Q, PS (S n) = SomeT _ Q /\ Red P Q) /\
 PS (S n) = NoneT _) /\
 (PS n = NoneT _ -> PS (S n) = NoneT _).

A *maximal reduction sequence* (or a *run*) is a reduction sequence whose last state is stable, or an infinite reduction sequence:

Definition `isMaxRedSeq` (PS:stateSeq) : Prop := isRedSeq PS /\
 (forall n P, PS n = SomeT _ P -> PS (S n) = NoneT _ -> Stable P).

One may observe that empty sequences are valid runs. In the encoding of formulas, we enforce the condition that a run starts with some state.

Encoding of Fairness

We formalize the notion of *strong fairness*. Informally, strong fairness says that any process that is infinitely often enabled is eventually reduced⁵. We need a few intermediate definitions. We say that P is a subprocess of Q when Q consists of the parallel composition of P with some other process(es). The predicate **reduced** $P \ Q \ R$ intuitively means Q reduces to R by reducing its subprocess P . The formal definition is omitted for lack of space.

We define what it means for a subprocess to be *enabled* and *eventually reduced*:

Definition `enabled` ($P : \text{proc}$) ($Q : \text{state}$) : $\text{Prop} :=$
`exists R, reduced P Q R.`

Definition `ev_reduced` ($P : \text{proc}$) ($PS : \text{stateSeq}$) : $\text{Prop} :=$
`exists n, (exists Q, (exists R,`
`PS n = SomeT _ Q /\ PS (S n) = SomeT _ R /\ reduced P Q R)).`

We define what it means for a property to hold *infinitely often*:

Definition `is_postfix` ($PS' \ PS : \text{stateSeq}$) : $\text{Prop} :=$
`exists n, (forall m, PS' m = PS (m + n)).`

Definition `infinitely_often` ($p : \text{state} \rightarrow \text{Prop}$) ($PS : \text{stateSeq}$) : $\text{Prop} :=$
`forall m, exists n, (exists P, PS (m + n) = SomeT _ P /\ p P).`

A *fair reduction sequence* is a state sequence such that there is no process that is infinitely often enabled but never reduced:

Definition `isFairRedSeq` ($PS : \text{stateSeq}$) : $\text{Prop} :=$
`forall PS', is_postfix PS' PS ->`
`forall P, infinitely_often (fun Q => enabled P Q) PS' ->`
`ev_reduced P PS'.`

3.2 Available Formulas

Our specification language consists of a set of logical and spatial formulas (of type **form**) and a set of temporal formulas (of type **tform**). The semantics of formulas is implemented by two satisfaction relations (**sat** of type **form** \rightarrow **state** \rightarrow **Prop** and **tsat** of type **tform** \rightarrow **state** \rightarrow **Prop**). The explicit distinction between logical and spatial formulas, and temporal formulas is required for the confluence properties introduced in the next section to hold. The informal semantics of basic formulas appears in Table 1. Observe that we make use of a predicate **Cong** that encodes the standard notion of *structural congruence* (which intuitively relates processes that only differ by spatial rearrangements).

⁵ *Weak fairness* says that a continuously and infinitely enabled process is eventually reduced. Strong fairness subsumes weak fairness.

Logical Formulas

<code>sat ISANY S</code>	iff	<code>True</code>
<code>sat NEG f S</code>	iff	<code>~ sat f S</code>
<code>sat OR f g S</code>	iff	<code>sat f S /\ sat g S</code>

Spatial Formulas

<code>sat INPUTS c f L#P</code>	iff	<code>Cong P (parP (inP c Q) R) and sat f L#(Q v) for any v</code>
<code>sat OUTPUTS c v f L#P</code>	iff	<code>Cong P (parP (outP c v Q) R) and sat f L#Q</code>
<code>sat CONSISTS f g L#P</code>	iff	<code>Cong P (parP Q R) with sat f L#Q and sat g L#R</code>

Temporal Formulas

<code>tsat (MAYEV f) S</code>	iff	for some run, there exists S' such that <code>Reds S S'</code> and <code>sat f S'</code>
<code>tsat (FMUSTEV f) S</code>	iff	for any fair run, there exists S' such that <code>Reds S S'</code> and <code>sat f S'</code>

Table 1
Basic Formulas

By way of example, we show the formal semantics of the **FMUSTEV** temporal formula. It is defined by quantification over all possible fair runs, as defined in the previous section:

```
Axiom FMUSTEV_satisfaction : forall P F, tsat (FMUSTEV F) P <->
(forall PS, PS 0 = SomeT _ P -> isMaxRedSeq PS -> isFairRedSeq PS ->
  exists Q, (exists n, PS n = SomeT _ Q /\ tsat F Q)).
```

In our implementation, the satisfaction relations are axiomatized. This is because the formula for negation does not respect the positivity constraints imposed by Coq. This problem has already been observed in [29]. This is not problematic as long as we do not study formally the properties of the formulas.

4 Collection of Lemmas

At this point, we are able to write a concurrent program P , a (temporal) property f , and we can try to prove `tsat f P` using Coq tactics. This direct approach is tedious because the Coq native tactics are too low-level and not adapted to the problem at hand. Our solution is to propose a collection of lemmas (and accompanying tactics) to facilitate formal proof.


The main difficulty in proving properties of concurrent programs is non-determinism. In order to prove a property for some program, one often needs to check all possible runs. This is at best costly and often impossible because there may be infinitely many runs or because some process is unknown. To deal with these

situations, we propose several lemmas based on confluence properties. In Sect. 4.1 we explain lemmas based on confluence properties and in Sect. 4.2 we give an overview of the whole library.

4.1 Confluence Properties

Basic Idea

We say that two reductions are *confluent* when they can be executed in either order to reach the same result. More precisely, if P is a process such that $P \xrightarrow{l_1} P_1$ and $P \xrightarrow{l_2} P_2$ are confluent, then for any P' such that $P_2 \xrightarrow{l_1} P'$, we have $P_1 \xrightarrow{l_2} P'$. Graphically, P has the following “diamond property”:



Since we know that, no matter the run, P necessarily reduces to P' , it is not always necessary to explore both runs to verify a **FMUSTEV** property. This is the basic idea behind lemmas based on confluence properties.

Partial Confluence and Linearized Channels

In order to identify “diamond properties”, we appeal to the notion of partial confluence, which is more general than confluence and often occurs in practice. We say that a reduction is *partially confluent* [21] when it is confluent with any other reduction. More precisely, if P is a process such that $P \xrightarrow{l_1} P_1$ is a partially confluent reduction, then for any reduction $P \xrightarrow{l_2} P_2$ and for any P' such that $P_2 \xrightarrow{l_1} P'$, we have $P_1 \xrightarrow{l_2} P'$.

The property of partial confluence is enjoyed by *linearized channels* [21]. A linearized channel is a generalization of a linear channel. It can be used more than once, but only in a sequential manner: an output process $\text{outP } c \vee P$ can reuse c again for output in P , an input process $\text{inP } c \wedge P$ can reuse c again for input in $P \vee$ for any \vee of the appropriate type. We introduce linearized channels in $\text{appl}\pi$ by adding some boolean information to the type of channels **chan** and by adding a new constructor to the type of processes **proc**:

Axiom `chan : Set -> bool -> Set.`

Inductive `proc : Type :=`

```
...
| nuP : forall A, (chan A false -> proc) -> proc (* non-linearized *)
| nuPl : forall A, (chan A true -> proc) -> proc. (* linearized *)
```

The operational semantics is modified accordingly.

For the time being, we assume that linearized channels are correctly annotated. Verification that a process is well-annotated can be done by the type system proposed in [20], Sect. 6.

Sample Confluence Property

The following example is taken from our library:

```
Axiom conf_red_com : forall L P P' A (c:chan A true),
  well_annotated (L#P) ->
  Redwith (L#P) (epsilon c) (L#P') ->
  forall f Lf, tfree_vars Lf f ->
  ~ in_Channels c Lf ->      (* f does not depend on the      *)
  forall K, guard K P' ->    (* channels that are consumed  *)
  inter Lf K nilC ->        (* or revealed by communication *)
  tsat (FMUSTEV f) (L#P') ->
  tsat (FMUSTEV f) (L#P).
```

Intuitively, it says that if $L\#P$ reduces to $L\#P'$ by a linearized communication, then in order to prove $\text{tsat (FMUSTEV } f) L\#P$, it is sufficient to prove $\text{tsat (FMUSTEV } f) L\#P'$, modulo some syntactical condition on f . Currently, these lemmas are axiomatized. They are similar to partial order reduction techniques used in model checking. See [4] for a pencil-and-paper evidence of their validity.

4.2 Library Overview

The library consists of the $\text{ap}\pi$ language as defined in Sect. 2 (extended with linearized channels), the specification language as defined in Sect. 3, and a collection of lemmas. Although the library is very large (at the time of this writing, 42 scripts, 20969 lines), only a few lemmas are axiomatized (most axioms have actually been discussed in this paper). See [3] for details.

Not all lemmas are equally important. During formal proof, the most important lemmas are those that simplify the goal. For example, confluence properties such as the one seen above are such lemmas: they basically act by simplifying the process that appears in the goal. Similarly, the properties of the formulas of the specification language (distributivity laws, etc.) act by simplifying the formula that appears in the goal.

There are a large number of lemmas that are not intended to be used directly during formal proof but that are very important because they are ubiquitously used to prove other lemmas. Such technical lemmas prove properties about the $\text{ap}\pi$ language (injection, inversion) whose proofs are not immediate because of our use of dependent types, and properties about structural congruence (e.g., structural congruence is a bisimulation).

5 Case Study

We evaluate feasibility of verification of concurrent programs using our library. Our case study is the SMTP receiver part of an existing mail server. In short, this program receives and processes SMTP commands, sends back SMTP replies and queues received electronic mail.

We chose this application for the purpose of comparison. Indeed, we have already performed verification of this application in Coq using a different approach that consists of building a faithful functional model [2]. In short, the original Java implementation was turned into a Coq function using monadic style programming, third-party programs (client and file-system) were modeled using Coq predicates and non-software aspects were modeled using functional constructs (for example, non-deterministic system failures were modeled using infinite lists to serve as test oracles). Arguably, this approach has little overhead because it takes advantage of the Coq built-in support for functional programs. Therefore, comparison should highlight the overhead of using our library for verification.

In the following, we first explain how we model the mail server using our library, and then we comment on the formal proof that it correctly implements the SMTP protocol. See [3] for details.

Modeling of the Main Program

The mail server is modeled as a process **work** that is itself the parallel composition of several processes that handle incoming SMTP requests: **get_helo_def**, etc. The state of the application is reified as a data structure that is communicated from one subprocess to the other. The flow of communication reproduces the flow of control of the Java program. Subprocesses correspond to the Java methods that (are supposed to) implement the SMTP protocol. The reified state corresponds to fields of the server object:

Definition work

```
(c1:InputStream) (c2:OutputStream) (tofs:ToFileSystem) : proc :=
  let st := initial_state in
  nuPl (fun heloc:chan STATE _ =>
    nuPl (fun mailc:chan STATE _ =>
      nuPl (fun rcptc:chan STATE _ =>
        parP (rinP heloc (get_helo_def heloc mailc))
          (parP (rinP mailc (get_mail_def mailc rcptc))
            (parP (rinP rcptc (get_rcpt_def mailc rcptc))
              (OutAtom heloc st)))))).
```

Modeling benefits from the fact that $\text{ap}\pi$ is based on the π -calculus. The connections between the mail server and third-party programs (client and file-system) are modeled using channels (instead of sockets in the original Java implementation) that are aggregated into the reified state and “move” around with the state during computation. Acknowledgments are modeled by a typical π -calculus idiom: a fresh channel is sent to receive the acknowledgment on it.

Modeling of Third-party Programs

Third-party programs are modeled by Coq predicates. For example, the client is modeled by predicates (`speaks_valid_protocol` and `acknowledges_replies`) that implement the SMTP protocol as defined in RFC 821.

Modeling of System Errors

System errors are modeled by channels. A system failure (resp. a network error) is modeled by outputting some value along the channel `system_failure_chan` (resp. `IOexn_chan`). Since non-determinism is inherent to $\text{ap}\pi$, we can model non-deterministic system failures by a process:

```
Definition may_fail := nuP (fun x => parP (InAtom x)
  (parP (OutAtom x tt)
    (inP x (fun _ => OutAtom system_failure_chan tt)))).
```

This is more elegant than infinite lists that serve as test oracles in the functional model discussed above. Indeed, the process `may_fail` is clearly separated from the model of the main program, so that we can extract an ML program for the server without pollution from the modeling of system errors.

Formal Proof

We have formally proved that the parallel composition of the mail server, a valid client, a valid file-system, and a non-deterministic system failures generator ends up with a successful termination (modeled by channel `result_chan`, similarly to system errors), a system failure or a network error (formula `reports_succ_or_error`):

```
Definition reports_succ_or_error : form :=
  OR (OUTPUTS result_chan tt ISANY)
  (OR (OUTPUTS IOexn_chan tt ISANY)
    (OUTPUTS system_failure_chan tt ISANY)).
```

```
Theorem server_accepts_valid_protocol :
  forall Client : InputStream -> OutputStream -> proc,
    (forall s y, (valid_client (Client s y))) ->
  forall file_system : ToFileSystem -> proc,
    (forall tofs, valid_fs tofs (file_system tofs)) ->
  (* channels for termination detection are distinct *)
  is_set (result_chan & IOexn_chan & system_failure_chan & nilC) ->
  tsat (FMUSTEV (STAT reports_succ_or_error))
  (result_chan & IOexn_chan & system_failure_chan & nilC #
  nuPl (fun s1:InputStream =>
  nuPl (fun s2:OutputStream =>
  nuPl (fun tofs:ToFileSystem =>
    parP (Client s1 s2) (parP (file_system tofs)
      (parP (work s1 s2 tofs) may_fail)))))).
```

Verification using our library requires 3927 commands. This is large compared to the 1059 commands required by the verification using the functional model. However, there are several ways to reduce the size of the proof. In particular, we used for verification a confluence property that is weaker (but easier to use in practice) than the one presented in Sect. 4.1. Also, it should be observed that the $\text{appl}\pi$ model is more satisfactory than the functional model in many respects: the **work** process takes multi-threading into account and it can easily be run as an ML program, which was not the case of the functional model.

6 Conclusion

In this paper, we proposed a Coq library to verify realistic concurrent programs. We have formalized a modeling language based on the π -calculus that is convenient to write and run (models of) realistic concurrent programs. We have introduced a specification language based on spatial logics extended with the notion of strong fairness. In order to facilitate formal proof, we have built a collection of lemmas among which confluence properties of the modeling language significantly simplify proofs. We have evaluated the feasibility of our approach by verifying a non-trivial application using our library.

Related Work

There exist several formalizations of pure concurrent calculi in proof assistants. In Coq, Hirschkoﬀ proposes a first-order abstract syntax encoding of the π -calculus and formalizes proof techniques [17]; Despeyroux formalizes a proof of subject reduction for the π -calculus [13]; Honsell et al. formalize the foundational paper on the π -calculus [18]; Scagnetto and Miculan formalize the ambient calculus (a derivative of the π -calculus) and its spatial logic [29]. In Isabelle, Röckl et al. propose a HOAS encoding of the π -calculus and formalize the Theory of Contexts [28]. This work focuses on the formalization of the theory of pure concurrent calculi. In contrast, we are concerned with verification of realistic concurrent programs and we aim at building a practical library for that purpose.

The verification of concurrent programs using proof assistants is also addressed using the UNITY formalism. In particular, there exist several formalizations of compositional reasoning [16,25] that is useful to tackle realistic examples. Our work is complementary: we use the π -calculus as an underlying formalism and therefore we can benefit from known analyses to formalize additional proof techniques (e.g., lemmas based on confluence properties).

Watkins et al. propose a logical framework [31] with built-in facilities for reasoning about concurrency. Using this concurrent logical framework (CLF), Cervesato et al. encode several concurrent systems [8], including the π -calculus. Although the authors have not addressed directly the issue of verification of realistic concurrent programs, it seems that an implementation of CLF may ease the development of a library similar to ours.

Coupet-Grimal [9] proposes an encoding of linear temporal logic in Coq. Temporal formulas are defined for an abstract transition system and their properties are collected into a library that has been used to prove correctness of a garbage-collection algorithm [10]. It would be useful to integrate similar reasoning on temporal formulas for our language.

Future Work

As stated in Sect. 4.1, we assume that channels are annotated so as to reflect partial confluence. In order to verify that channels are correctly annotated, we plan to formalize an adequate type system inside Coq, similarly to the work by Gay [15] who formalizes the type system of Kobayashi et al. [21] in Isabelle. We also plan to provide mechanical proofs for the lemmas based on confluence properties that are axiomatized for the time being.

We have been formalizing new formulas (such as fixed points) to enhance expressiveness but they are not yet integrated in the library. In order to reduce the size of formal proofs, we are improving automation and investigating additional proof techniques based on other type systems, such as Kobayashi's type system for lock-freedom [19].

References

- [1] Abadi, Martín, and Cédric Fournet, *Mobile values, new names, and secure communication*, in *28th ACM Symposium on Principles of Programming Languages (POPL 2001)*, p. 104–115, ACM Press.
- [2] Affeldt, Reynald, and Naoki Kobayashi, *Formalization and verification of a mail server in Coq*, in *International Symposium on Software Security*, volume 2609 of *Lecture Notes in Computer Science*, p. 217–233, Springer-Verlag, Feb. 2003.
- [3] Affeldt, Reynald, and Naoki Kobayashi, *A Coq Library for Verification of Concurrent Programs*, <http://staff.aist.go.jp/reynald.affeldt/applpi>, Coq scripts.
- [4] Affeldt, Reynald, and Naoki Kobayashi, *Partial Order Reduction for Verification of Spatial Properties of Pi-Calculus Processes*, in *11th International Workshop on Expressiveness in Concurrency (EXPRESS 2004)*, *Electronic Notes in Theoretical Computer Science*, 128(2):151–168, Elsevier, 2005.
- [5] Black, Paul E., “Axiomatic Semantics Verification of a Secure Web Server,” Ph.D. thesis, Department of Computer Science, Brigham Young University, 1998.
- [6] Boulton, Richard, Andrew Gordon, Mike Gordon, John Harrison, John Herbert, and John Van Tassel, *Experience with embedding hardware description languages in HOL*, in *IFIP TC10/WG 10.2 International Conference on Theorem Provers in Circuit Design: Theory, Practice and Experience*, volume A-10 of IFIP Transactions, p. 129–156, Elsevier, 1992.
- [7] Cardelli, Luca, and Andrew D. Gordon, *Anytime, anywhere: modal logics for mobile ambients*, in *27th ACM Symposium on Principles of Programming Languages (POPL 2000)*, p. 365–377, ACM Press.
- [8] Cervesato, Iliano, Frank Pfenning, David Walker, and Kevin Watkins, *A concurrent logical framework II: Examples and applications*, Technical Report CMU-CS-02-102, Department of Computer Science, Carnegie Mellon University, Pittsburgh, PA, May 2003.
- [9] Coupet-Grimal, Solange, *An axiomatization of linear temporal logic in the calculus of inductive constructions*, *Journal of Logic and Computation*, 13(6):801–813, 2003.
- [10] Coupet-Grimal, Solange, and Catherine Nouvet, *Formal verification of an incremental garbage collector*, *Journal of Logic and Computation*, 13(6):815–833, 2003.
- [11] Dam, Mads, *Proof Systems for the pi-calculus Logics*, in “Logic for Concurrency and Synchronisation,” *Trends in Logic, Logica Library*, Kluwer, 2003.

- [12] Despeyroux, Joëlle, Amy Felty, and André Hirschowitz, *Higher-Order Abstract Syntax in Coq*, in *2nd International Conference on Typed Lambda Calculi and Applications (TLCA 1995)*, volume 905 of *Lecture Notes in Computer Science*, p. 124–138, Springer-Verlag, Apr. 1995.
- [13] Despeyroux, Joëlle, *A higher-order specification of the π -calculus*, in *IFIP Conference on Theoretical Computer Science 2000*, volume 1872 of *Lecture Notes in Computer Science*, p. 425–439, Springer-Verlag, Aug. 2000.
- [14] Filliâtre, Jean-Christophe, *Why: a multi-language multi-prover verification tool*, Research Report 1366, LRI, Université Paris Sud, Mar. 2003.
- [15] Gay, Simon J., *A framework for the formalisation of pi calculus type systems in Isabelle/HOL*, in *Theorem Proving in Higher Order Logics (TPHOLs 2001)*, volume 2152 of *Lecture Notes in Computer Science*, p. 217–232, Springer-Verlag, Sep. 2001.
- [16] Heyd, Barbara, and Pierre Crégut, *A modular coding of UNITY in Coq*, in *Theorem Proving in Higher Order Logics (TPHOLs 1996)*, volume 1125 of *Lecture Notes in Computer Science*, p. 251–266, Springer-Verlag, Aug. 1996.
- [17] Hirschhoff, Daniel, “Mise en œuvre de preuves de bisimulation,” Ph.D. thesis, École Nationale des Ponts et Chaussées, 1999.
- [18] Honsell, Furio, Marino Miculan, and Ivan Scagnetto, *π -calculus in (co)inductive type theory*, *Theoretical Computer Science*, 253(2):239–285, Feb. 2001.
- [19] Kobayashi, Naoki, *A type system for lock-free processes*, *Information and Computation*, 177(2):122–159, Sep. 2002.
- [20] Kobayashi, Naoki, *Type systems for concurrent programs*, Tutorial, in *UNU/IIST 10th Anniversary Colloquium, March 2002, Lisbon, Portugal*, Springer-Verlag, 2002.
- [21] Kobayashi, Naoki, Benjamin C. Pierce, and David N. Turner, *Linearity and the Pi-Calculus*, in *23rd ACM Symposium on Principles of Programming Languages (POPL 1996)*, p. 358–371, ACM Press.
- [22] Manolios, Panagiotis, “Mechanical Verification of Reactive Systems,” Ph.D. thesis, The University of Texas at Austin, Department of Computer Sciences, Austin, TX, 2001.
- [23] Melham, Thomas F., *A Mechanized Theory of the π -calculus in HOL*, *Nordic Journal of Computing*, 1(1):50–76, 1995.
- [24] Milner, Robin, Joachim Parrow, and David Walker, *A calculus of mobile processes, parts I and II*, *Information and Computation*, 100(1):1–77, 1992.
- [25] Paulson, Lawrence C., *Mechanizing a theory of program composition for UNITY*, *ACM Transactions on Programming Languages and Systems*, 23(5):626–656, 2001.
- [26] Pierce, Benjamin C., and David N. Turner, *Pict: A programming language based on the pi-calculus*, in “Proof, Language and Interaction: Essays in Honour of Robin Milner,” MIT Press, 2000.
- [27] Pierce, Benjamin C., and Jérôme Vouillon, *Specifying a file synchronizer*, Draft, Mar. 2002.
- [28] Röckl, Christine, Daniel Hirschhoff, and Stefan Berghofer, *Higher-order abstract syntax with induction in Isabelle/HOL: Formalizing the pi-calculus and mechanizing the theory of contexts*, in *Foundations of Software Science and Computation Structures (FOSSACS 2001)*, number 2030 in *Lecture Notes in Computer Science*, Springer-Verlag, Apr. 2001.
- [29] Scagnetto, Ivan and Marino Miculan, *Ambient calculus and its logic in the calculus of inductive constructions*, in *3rd International Workshop on Logical Frameworks and Meta-Languages (LFM 2002)*, *Electronic Notes in Theoretical Computer Science*, 70(2), Elsevier, 2002.
- [30] The Coq Development Team, “The Coq Proof Assistant Reference Manual,” INRIA, 2004.
- [31] Watkins, Kevin, Iliano Cervesato, Frank Pfenning, and David Walker, *A concurrent logical framework I: Judgments and properties*, Technical Report CMU-CS-02-101, Department of Computer Science, Carnegie Mellon University, Pittsburgh, PA, May 2003.
- [32] Yu, Shen-Wei, “Formal Verification of Concurrent Programs Based on Type Theory,” Ph.D. thesis, Department of Computer Science, University of Durham, 1998.